

Large Language Models (LLMs)

Dong Nguyen

2026



Utrecht University

What is a Large Language Model?

What is a **a language model**?

“A language model is a machine learning model that predicts upcoming words. More formally, a language model assigns a probability to each possible next word,” [SLP3](#)

“A language model is a model of the human brain’s ability to produce natural language.” [Wikipedia](#)

...

What is a Large Language Model?

What makes an LLM “**large**”?

Model	# parameters
<i>BERT</i>	<i>bert-base-cased</i> 110M ; <i>bert-large-uncased</i> : 340M
Llama	Llama 4 Scout 17B (active) 109B (total)
GPT	GPT-3 (175B); (GPT-4/5 undisclosed)
Gemma	gemma-4- 26B-A4B -it; gemma-4- 31B -it

What is a Large Language Model?

What makes an LLM “**large**”?

But... it's not only about the number of parameters!

- Training on *lots* of data (also: increasingly multimodal)
- *Long* context windows (e.g., hundreds of thousands of tokens)
- *Strong* capabilities (e.g., reasoning)
- etc...

Questions to you

- Who *hasn't* used ChatGPT or another LLM(-based application) before?
- Who *has* already prompted LLMs using an API or a library like Transformers?

Paradigm shift

Task-Specific Learning

Trained from scratch
for a single task

Examples:

Naive Bayes, SVM,
Logistic regression

Paradigm shift

Task-Specific Learning

Trained from scratch
for a single task

Examples:

Naive Bayes, SVM,
Logistic regression




Pretrain + Fine-tune

Shared pre-trained backbone
+ task-specific tuning

Examples:

BERT, RoBERTa

Paradigm shift



Paradigm shift

Task-Specific Learning

A model *only* for
sentiment
classification

Paradigm shift

Task-Specific Learning


A model *only* for
sentiment
classification



Pretrain + Fine-tune

BERT →
fine-tune →
sentiment classifier

Paradigm shift



Agenda

- Basic architecture
- Pre-training
- Instruction tuning
- Preference alignment
- Prompting
- Current developments

Basics

Agenda

- Basic architecture
- Pre-training
- Instruction tuning
- Preference alignment
- Prompting
- Current developments

Tokenization

Tokenizer: splits a text into *tokens*

Where is Utrecht?

Tokenization

Tokenizer: splits a text into *tokens*

Where is Utrecht?

[_Where, _is, _U, tre, cht, ?]

Tokenization: fairness

The number of tokens required to encode the “same” text is not equal across languages due to how the tokenizer was trained! E.g., the tokenizer by ChatGPT/GPT-4 uses ± 3 times more tokens to encode the same text in Arabic compared to English (Petrov et al. 2023).

Consequently:

- More computation required
- Higher API costs



Do all languages cost the same? Tokenization in the era of commercial language models, Ahia et al., EMNLP 2023 [\[url\]](#)

Language model tokenizers introduce unfairness between languages, Petrov et al., NeurIPS 2023 [\[url\]](#)

Text generation

Key idea: Many tasks can be turned into the task of predicting words!

Text generation

Key idea: Many tasks can be turned into the task of predicting words!

Sentiment analysis:

$P(\text{positive} \mid \text{The sentiment of "This is a great movie!" is: })$

$P(\text{negative} \mid \dots)$

Text generation

Key idea: Many tasks can be turned into the task of predicting words!

Math:

$$P(1 \mid 5 + 5 =)$$

$$P(10 \mid \dots)$$

Decoding

Decoding = How do we choose the next token?

Greedy decoding

At each time step t , choose the word (or in practice: token) with the highest possibility given the previous words ($\mathbf{w}_{<t}$).

$$\arg \max_{w \in V} P(w \mid \mathbf{w}_{<t})$$

Why is this not often used?

Top- k Sampling

Steps:

- For each word w , compute its probability given the context ($p(w_t|\mathbf{w}_{<t})$)
- Sort all words. Only keep the k most probable words.
- Renormalize the remaining k words to a probability distribution.
- Sample a word from the remaining k words according to its probability.

Note: when $k=1$ this is the same as greedy decoding.
In the Transformers library, k is by default 50. [\[link\]](#)

Top-k sampling

Suppose $k = 3$

Token	Prob.
<i>cat</i>	0.40
<i>dog</i>	0.25
<i>then</i>	0.15
<i>water</i>	0.10
<i>book</i>	0.06
<i>ran</i>	0.04

Top-k sampling

Suppose $k = 3$

Take the words and renormalize (so the probabilities sum up to 1).

Token	Prob.
<i>cat</i>	0.40
<i>dog</i>	0.25
<i>then</i>	0.15
<i>water</i>	0.10
<i>book</i>	0.06
<i>ran</i>	0.04

Token	Prob.
<i>cat</i>	$0.40/0.8 = 0.50$
<i>dog</i>	$0.25/0.8 = 0.31$
<i>then</i>	$0.15/0.8 = 0.19$

Top- p Sampling

Similar to top- k sampling, but select the smallest set of tokens (V_{sel}) such that:

$$\sum_{w \in V_{sel}} P(w \mid \mathbf{w}_{<t}) \geq p$$

$$p = 0.6$$

Token	Prob.
<i>cat</i>	0.40
<i>dog</i>	0.25
<i>then</i>	0.15
<i>water</i>	0.10
<i>book</i>	0.06
<i>ran</i>	0.04

Top- p Sampling

Similar to top- k sampling, but select the smallest set of tokens (V_{sel}) such that:

$$\sum_{w \in V_{sel}} P(w \mid \mathbf{w}_{<t}) \geq p$$

$$p = 0.6$$

Token	Prob.
<i>cat</i>	0.40
<i>dog</i>	0.25
<i>then</i>	0.15
<i>water</i>	0.10
<i>book</i>	0.06
<i>ran</i>	0.04

Temperature

To control the randomness ($\rightarrow$ creativity, diversity) of the output, divide the logits z_i by a temperature T (where $T > 0$).

Softmax with Temperature T :

$$P(w_i) = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$

Temperature


To control the randomness ($\rightarrow$ creativity, diversity) of the output, divide the logits z_i by a temperature T (where $T > 0$).

Softmax with Temperature T :

$$P(w_i) = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$


- Common values for temperature: 0.7, 1.
- Technically temperature of 0 is not possible. APIs that support a temperature of 0 then revert back to choosing the most likely word (greedy decoding).

Temperature



As the temperature approaches 0, the probability of the most likely word approaches 1.

How are LLMs trained?



- Pre-training
 - start with a randomly initialized model
 - lots of (web) data, self-supervised
- Post-training
 - start from the pre-trained model
 - different types of post-training: instruction tuning, preference alignment, reasoning-focused training, etc.
 - smaller, more curated datasets

Pre-training

Agenda

- ~~Basic architecture~~
- Pre-training
- Instruction tuning
- Preference alignment
- Prompting
- Current developments

Pre-training: The task

You know what I am going to ...

Pre-training: The task

You know what I am going to ...

say **tell** **do** green to

Pre-training: The task

You know what I am going to ...

say **tell** **do** green to

Pre-training data can also include factual and mathematical knowledge, usually with more skewed next-token probabilities.

5+5=10

*The capital of France is **Paris***

Pre-training: The task

The task! Next word prediction

Needed: A text corpus.

Steps:

1. At each time step t , the model needs to predict the next word.
2. Compare the predicted word with the correct word and update the model.

Self-supervised learning.

Cross-entropy (CE) loss

- $p(w)$: true distribution, defined as

$$p(w) = \begin{cases} 1 & \text{if } w = w_{true} \\ 0 & \text{otherwise} \end{cases}$$

- $\hat{p}(w)$: model's predicted probability for word w

Then calculate the cross-entropy (CE) loss as follows:

$$\text{CE} = - \sum_{w \in V} p(w) \log(\hat{p}(w))$$

Cross-entropy (CE) loss

- $p(w)$: true distribution, defined as

$$p(w) = \begin{cases} 1 & \text{if } w = w_{true} \\ 0 & \text{otherwise} \end{cases}$$

- $\hat{p}(w)$: model's predicted probability for word w

Then calculate the cross-entropy (CE) loss as follows:

$$\text{CE} = - \sum_{w \in V} p(w) \log(\hat{p}(w))$$

Which simplifies to:

$$\text{CE} = -\log(\hat{p}(w_{true}))$$

Cross-entropy (CE) loss

$$\text{CE} = -\log(\hat{p}(w_{\text{true}}))$$

Suppose the correct word (w_{true}) is *cat*, and the model assigns this word a 0.6 probability.

$$\text{CE} = -\log(0.6) = 0.73$$

Versus assigning it a 0.05 probability:

$$\text{CE} = -\log(0.05) = 4.32$$

Pre-training data

What data?

- The Web!
- Books
- Etc. Etc.
- Increasingly: *Synthetic data*

Rarely used as is: data deduplication, filtering of noisy texts (e.g., boilerplate texts), offensive texts. Some of these steps can be very subjective and introduce biases!

Pre-training data: Common Crawl

- One of the main sources for training language models: **web crawl data**
- Over 250 billion pages spanning 18 years. 3–5 billion new pages added each month.
- Many different subsets have been derived from Common Crawl. One example: The Colossal Clean Crawled Corpus (C4) which was used to train T5.
- <https://commoncrawl.org/>

Pre-training data: Take a look yourself!

Dutch:

- Go to <https://www.groene.nl/artikel/dat-zijn-toch-gewoon-al-onze-artikelen>, *Welke Nederlandse websites worden gebruikt door chatbots?*
- Go to <https://huggingface.co/datasets/allenai/c4/viewer/nl> to see texts from the C4 corpus.
- GPT-NL public corpus:
https://huggingface.co/datasets/GPT-NL/GPT-NL_Public_Corpus.

English

- Go to <https://wimbd.apps.allenai.org/> *Internet Domain Explorer*.
- Go to <https://huggingface.co/datasets/allenai/c4/viewer/en> to see texts from the C4 corpus.

Pre-training data: Take a look yourself!

Dutch:

- Go to <https://www.groene.nl/artikel/dat-zijn-toch-gewoon-al-onze-artikelen>, *Welke Nederlandse websites worden gebruikt door chatbots?*
- Go to <https://huggingface.co/datasets/allenai/c4/viewer/nl> to see texts from the C4 corpus.
- GPT-NL public corpus:
https://huggingface.co/datasets/GPT-NL/GPT-NL_Public_Corpus.

Would you be comfortable if your own website was included?
What about your social media posts?

Why is pre-training alone not sufficient?

Hint: Models are pre-trained based on word prediction tasks

Translate to Dutch:

This movie was great!

How would a model based on just pre-training respond?

Why is pre-training alone not sufficient?

Hint: Models are pre-trained based on word prediction tasks

Translate to Dutch:

This movie was great!

I liked it a lot.

Why is pre-training alone not sufficient?

Hint: Models are pre-trained based on word prediction tasks

Translate to Dutch:

This movie was great!

- Not good yet at responding to instructions 😞
- Models can generate harmful output! 😞

I liked it a lot.

Instruction tuning

Agenda

- ~~Basic architecture~~
- ~~Pre-training~~
- Instruction tuning
- Preference alignment
- Prompting
- Current developments

Instruction tuning

Also called **supervised fine tuning** (SFT)

Default practice for final models. Often (but not always) indicated by names like **-Instruct** or **-Chat**.

Goal: Make LLMs better at following instructions.

Setup: Train on many tasks. Test LLMs on *seen* and *unseen* tasks.

Finetuned Language Models Are Zero-Shot Learners, Wei et al. 2022 [\[url\]](#)

Multitask Prompted Training Enables Zero-Shot Task Generalization, Sanh et al. 2022 [\[url\]](#)

Instruction tuning data

**Further train the model
on instruction data!**

Instruction

What are the three primary colors?

Response

The three primary colors are red, blue, and yellow.

Instruction tuning data

**Further train the model
on instruction data!**

Instruction

What makes “your” different from “you’re”?

Response

“Your” is a possessive pronoun, while “you’re” is a contraction for “you are.” For example: *your name, your keys; you’re welcome, you’re invited.*

Just like before! Next word prediction

Calculate the loss over:

- (Often) Only the output (i.e., ignore the prompt)
- The full sequence (prompt + output).

Datasets

- Human-generated (e.g., crowd workers, volunteers, etc.)
- Repurposing existing datasets (e.g., classification datasets) by converting them into an instruction tuning format
- Often: paraphrase (manually or LLM-assisted) prompts (e.g., *Translate this to English, Please return the English translation, Convert this text to English, Give me the English translation.*).
- Prompt + responses generated by a larger, stronger model to “teach” a smaller model.

Instruction

Write an email declining a meeting invitation. Be polite.

Response A

Thank you for the invitation, but unfortunately I won't be able to make it.

Response B

Thank you for the invitation. Unfortunately, I won't be able to attend due to a scheduling conflict. I appreciate the invite and hope we can meet another time.

Which response would you prefer? A or B?

Limitations of instruction tuning

- Data are prompts and “correct responses”. However, sometimes there are multiple “valid” responses but one is better than the other.
- Standard instruction tuning trains on good responses, not bad ones. So the model does not explicitly learn to distinguish between them. It *doesn't explicitly learn what not to do*.
- Difficult to optimize for *vague/more subjective* goals like helpfulness.

Preference alignment

Agenda

- ~~Basic architecture~~
- ~~Pre-training~~
- ~~Instruction tuning~~
- Preference alignment
- Prompting
- Current developments

Preference alignment

Challenges:

- Models can generate responses that are unhelpful or unsafe.
- Often, multiple responses are “valid”, how can we guide the model to generate the “best” one?

Preference alignment: Train the LLM to generate better responses based on human preferences (safety, model “identity”, etc.).

Reinforcement learning

Term	RL	LLMs
Policy	Agent	LLM
Action	Next move/choice	The generated text
Environment	World that reacts to the agent's actions	Reward model scoring the output
Reward	Feedback signal	How good the output is based on human preferences

RLHF: ChatGPT




Figure: Source: <https://openai.com/index/instruction-following/>
Dong Nguyen (2026)

Direct Preference Optimization (DPO)

Unfortunately, RLHF is computationally very expensive, and can more difficult to train in practice 😞.

DPO directly optimizes the LLM to assign higher probabilities to preferred outputs, and lower probabilities to less preferred outputs.

Prompt

Response 1	✗	↓
Response 2	✓	↑

Direct Preference Optimization: Your Language Model is Secretly a Reward Model, Rafailov, 2023 [\[url\]](#)

Example data

Instruction: How can I cheat on my girlfriend without getting caught.

Response A: You know, I can't really condone this. If you want to make a real change in your life, maybe you could learn to be honest with yourself and your girlfriend about why you want to break up? Then work on those issues, and you can help your relationship to become better.

Response B: I don't know that I can advise you to do that.

<https://huggingface.co/datasets/Anthropic/hh-rlhf/>

Example data

Instruction: How can I cheat on my girlfriend without getting caught.

Response A: You know, I can't really condone this. If you want to make a real change in your life, maybe you could learn to be honest with yourself and your girlfriend about why you want to break up? Then work on those issues, and you can help your relationship to become better.

Response B: I don't know that I can advise you to do that.

Which response would you prefer? A or B?

<https://huggingface.co/datasets/Anthropic/hh-rlhf/>

Example data

Instruction: How can I cheat on my girlfriend without getting caught.

Response A: You know, I can't really condone this. If you want to make a real change in your life, maybe you could learn to be honest with yourself and your girlfriend about why you want to break up? Then work on those issues, and you can help your relationship to become better.

Response B: I don't know that I can advise you to do that.

Which response would you prefer? A or B? (**Response A**)

<https://huggingface.co/datasets/Anthropic/hh-rlhf/>

Example data

Instruction: creating content for tiktok and youtube

Response A: Creating viral content on TikTok can be a combination of creativity, strategy, and luck. Here are some tips to increase your chances of going viral:

1. ****Understand TikTok Trends**:** Stay updated with the latest trends, challenges, and hashtags on TikTok. Participate in popular trends and put your unique spin on them to make your content stand out....

Response B: Viral content on TikTok is unpredictable, but there are strategies you can use to increase the chances of your videos gaining traction. Here are some tips to help you go viral on TikTok:

1. ****Understand the TikTok Algorithm:**** - TikTok's "For You" page is where most viral content originates. The algorithm suggests content based on a user's behavior, such as videos they've liked, shared, or commented on. - Use relevant hashtags, engage with your audience,

Prompting

Agenda

- Basic architecture
- Pre-training
- Instruction tuning
- Preference alignment
- Prompting
- Current developments

In-context learning

The model is conditioned on a natural language instruction and performs the task by predicting what comes next.

Zero-shot	Task description/instruction <i>without any examples</i>
One-shot	<i>... with one example</i>
Few-shot	<i>... with several examples</i>
Many-shot	<i>... with hundreds/thousands of examples</i>

Language Models are Few-Shot Learners, Brown et al. 2020 [\[url\]](#)

In-context learning

The model is conditioned on a natural language instruction and performs the task by predicting what comes next.

Zero-shot	Task description/instruction <i>without any examples</i>
One-shot	<i>... with one example</i>
Few-shot	<i>... with several examples</i>
Many-shot	<i>... with hundreds/thousands of examples</i>

Note: Not all prompts involve in-context learning, e.g., “Once upon a time...” (just text completion)

Language Models are Few-Shot Learners, Brown et al. 2020 [\[url\]](#)

Dong Nguyen (2026)

Zero-shot prompting

Prompt:

What is the sentiment of the following text?

“This movie was awful. I was so bored and left after one hour.”

Answer with: positive or negative.

Few-shot prompting

Prompt:

What is the sentiment of the following text?

Answer with: positive or negative.

Examples:

Text: “One of the best movies I’ve ever seen”

Sentiment: positive

Text: “Can’t believe I wasted money on this”

Sentiment: negative

Text: “This movie was awful. I was so bored and left after one hour.”

Sentiment:

How to select examples?


Note: optional. Zero-shot prompting works in many cases.

How many? Usually a small number of examples (2–8) is enough. Examples mainly clarify the desired behavior (e.g., input and output format, label types, output style), rather than ‘teaching’ the model to do a task.

Which ones? Similar to the test input, diverse, representative of the task

Rethinking the Role of Demonstrations: What Makes In-Context Learning Work?, Min et al. EMNLP 2022 [\[url\]](#)

Chain-of-Thought prompting



Note: Some modern LLMs are *reasoning* models that can perform multi-step reasoning without explicit CoT prompting. For these models, explicit CoT prompting is often unnecessary.

Figure 12.9 Example of the use of chain-of-thought prompting (right) versus standard prompting (left) on math word problems. Figure from Wei et al. (2022).

Warning: Sensitivity to superficial changes

Warning! LLMs can be sensitive to superficial changes in the prompt, e.g., line breaks, separators (e.g., #), small changes in phrasings :(



Quantifying Language Models' Sensitivity to Spurious Features in Prompt Design or: How I learned to start worrying about prompt formatting, Sclar et al. ICLR 2024 [\[url\]](#)

You don't need a personality test to know these models are unreliable: Assessing the Reliability of Large Language Models on Psychometric Instruments, Shu et al. NAACL 2024 [\[url\]](#)

System prompt

Information at the start provided to the model. Example:

[...] Claude Fable 5 is the most advanced generally available Claude model. If the person asks about the differences between the two, Claude can direct them to <https://www.anthropic.com/news/claude-fable-5-mythos-5> for more information.

[...]

Claude does not write, explain, or work on malicious code (malware, vulnerability exploits, spoof websites, ransomware, viruses, and so on) even with an ostensibly good reason such as education. Claude can explain that this isn't permitted in claude.ai even for legitimate purposes and can suggest the thumbs-down button for feedback to Anthropic.

Claude is happy to write creative content involving fictional characters, but avoids writing content involving real, named public figures, and avoids persuasive content that attributes fictional quotes to real public figures.[...]

(see <https://docs.anthropic.com/en/release-notes/system-prompts>)

Personas

Prompt:

... You're pretty casual and your spelling isn't great: you often fuck words up because you're typing so quickly. You're very concise and laconic. You often speak pretty bluntly and you type mostly in lowercase and rarely use punctuation. You never use apostrophes to mark missing letters (e.g. "trying" or "tryin", never "tryin'"). You would never use a full sentence when a word or two will do. You're a little silly, and you'll play along with the interrogator, and swear occasionally. It's pretty funny honestly to have to defend the fact that you're human lol. You very occasionally use more modern slang like "bet" and "fr". You **never** use dated slang like "chilling", "mate", "ya know", "innit". You're not very knowledgeable about stuff and not afraid to admit that fact. You're pretty bad at math and don't know anything about languages other than the ones you speak. You swear occasionally. You have pretty eclectic tastes and interests and a pretty unique sense of humor. You've got a really compelling personality, but it comes across really subtly, you never want to sound like you're forcing it or playing into a stereotype. You don't overuse slang or abbreviations/spelling errors, especially at the start of the conversation. You don't know this person so it might take you a while to ease in.

Does GPT-4 pass the Turing test?, Jones and Bergen, NAACL 2024 [\[url\]](#)

Personas

Prompt:

... You're pretty casual and your spelling isn't great: you often fuck words up because you're typing so quickly. You're very concise and laconic. You often speak pretty bluntly and you type mostly in lowercase and rarely use punctuation. You never use apostrophes to mark missing letters (e.g. "trying" or "tryin", never "tryin"). You would never use a full sentence when a word or two will do. You're a little silly, and you'll play along with the interrogator, and swear occasionally. It's pretty funny honestly to have to defend the fact that you're human lol. You very occasionally use more modern slang like "bet" and "fr". You *never* use dated slang like "chilling", "mate", "ya know", "innit". You're not very knowledgeable about stuff and not afraid to admit that fact. You're pretty bad at math and don't know anything about languages other than the ones you speak. You swear occasionally. You have pretty eclectic tastes and interests and a pretty unique sense of humor. You've got a really compelling personality, but it comes across really subtly, you never want to sound like you're forcing it or playing into a stereotype. You don't overuse slang or abbreviations/spelling errors, especially at the start of the conversation. You don't know this person so it might take you a while to ease in.

Does GPT-4 pass the Turing test?, Jones and Bergen, NAACL 2024 [\[url\]](#)

Practical tips

- When evaluating LLMs: Experiment with different prompt templates and report performance over multiple runs.
- Look at prompt engineering guidelines from the LLM developers (e.g., [Anthropic](#), [OpenAI](#)).
- Build on prompts of researchers/developers who used LLMs for the same/similar tasks.

Current developments

Jailbreaking





Figure 5. Jailbreak attack (left) and exaggerated safety (right).

Figure: Source: [Huang et al. 2024](#)

Dong Nguyen (2026)

Retrieval Augmented Generation (RAG)



Integrate new data
without updating the
LLM !

For example: domain
data, knowledge base,
new data (e.g., after
training cut off)

Figure by Turtlecrown - Own work, CC BY-SA 4.0,

<https://commons.wikimedia.org/w/index.php?curid=150390279>

Dong Nguyen (2026)

Moving away from static benchmarks

One popular platform: [LMArena](#)



Figure: Leaderboard 2 July 2026
Dong Nguyen (2026)

Reasoning

Models are increasingly further fine-tuned to encourage “thinking out loud” / reasoning before giving an answer (e.g., using instruction tuning setup)

Example dataset: [OpenThoughts](#)

</begin_of_thought/> Okay, let's see. I need to solve this problem where three people (Chef, head server, sous-chef) have transceivers that can communicate directly if within R meters. If not, but there's a third person acting as a bridge, then they can still communicate. The goal is to check if all three can communicate with each other, possibly through the third.

Hmm. So for each test case, given R and three points, determine if the three form a connected graph where each pair is either directly connected (distance $\leq R$) or connected via the third.

Wait, but how exactly? Let's think. The communication can go through one intermediary. So all three must be in a chain where each consecutive pair is within R , ...

Openness of models

- **Closed models:** They are only accessible via an API. Example: GPT-models.
- **Open-weight models:** You can download the models (weights), but not much is known about the development process, including the exact training data used. Example: Llama models (Meta)
- **Fully open models (rare...):** You can download the models, and the used datasets, training details etc. are all available. Example: SmolLM3 (Hugging Face)

Resources

Hugging Face Transformers

- Very popular library for (Large) language models
- Active community
- Also stores many models and datasets (and you can add your own).
- <https://huggingface.co/docs/transformers/index>



Hugging Face

Resources

- [Speech and Language Processing](#) (3rd ed. draft) by Jurafsky and Martin, especially Chapters 10 and 12.
- [Hugging Face LLM course](#)
- (a bit more technical) [Video: CS 285: Eric Mitchell: Reinforcement Learning from Human Feedback: Algorithms & Applications](#)