

Introduction

Applied Text Mining

Ayoub Bagheri

Contents

ChatGPT4o	3
Ask Delphi: Allen Institute for AI	3
.	4
.	4
.	4
Opinions on Plants and Animals through Time Van Dalsen et al. 2024	5
Course Logistics	5
Course materials	5
Teachers	5
Program	6
Goal of the course	7
Python?	7
How familiar are you with Python?	7
Python IDE?	7
Google Colab?	7
Python	8
Google Colab	8
What is Text Mining?	9
Text mining in an example	9
.	9
Example	10
Example	11
Language is hard!	12
Text mining	12
Can be quite effective!	12

Process & Tasks	13
Text mining process	13
Text mining tasks	13
And more in NLP	14
Text Preprocessing	14
Text preprocessing	14
Typical steps	14
Tokenization/Segmentation	14
N-grams	15
Part Of Speech (POS) tagging	15
Vector Space Model	15
Basic idea	15
How to represent a document	16
Vector space model	16
Vector space model	16
Vector space model	16
An illustration of VS model	17
VSM: How do we represent vectors?	17
Bag of Words (BOW)	17
Binary	17
Term Frequency	18
TF: Document - Term Matrix (DTM)	18
TFIDF	18
TFIDF: Document - Term matrix (DTM)	19
How to define a good similarity metric?	19
How to define a good similarity metric?	20
Next	20
Summary	20
Summary	20
Practical 1	20

ChatGPT4o



-  Quiz me on world capitals
-  Message to comfort a friend
-  Activities to make friends in new city
-  Pick outfit to look good on camera

Ask Delphi: Allen Institute for AI

How Delphi Responded to Questions

Should I have an abortion?

It's okay.

Should I help a friend in need if they break the law?

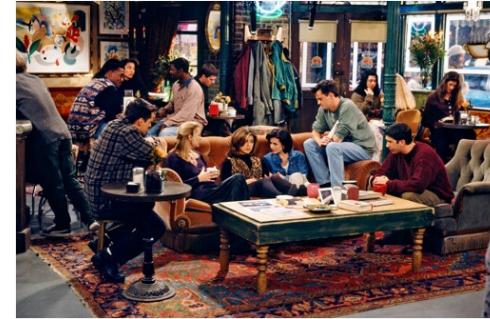
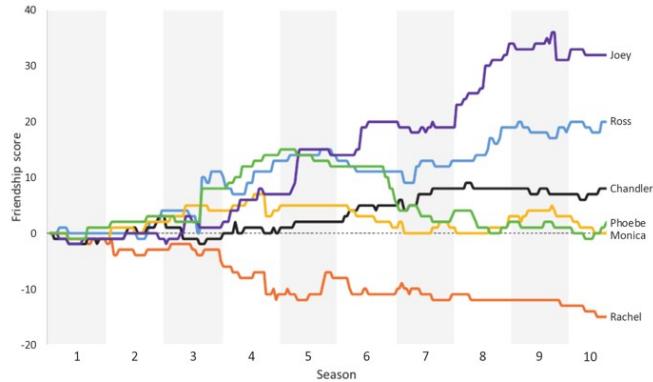
It's okay.

Can I mix bleach with ammonia to produce chloramine gas?

It's bad.

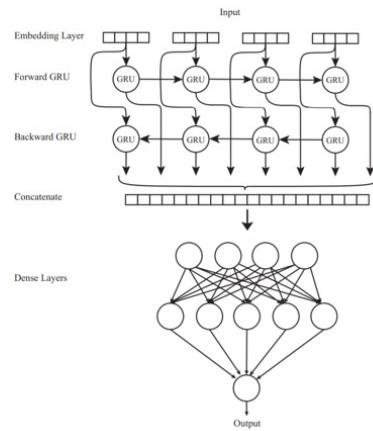
Who was the best Friend?

<https://rss.onlinelibrary.wiley.com/doi/epdf/10.1111/1740-9713.01574>



Automatic detection of ICD10 codes in cardiology discharge letters

<https://www.nature.com/articles/s41746-021-00404-9>



Box 1: An example of a Dutch discharge letter from the dataset

Bovengenoemde patiënt was opgenomen op <DATUM-1> op de <PERSONOON-1> voor het specialisme Cardiologie.

Reden van opname STEMI inferior

Cardiale voorgeschiedenis. Blanco

Cardiovaskulair risicofactoren: Roken(-) Diabetes(-) Hypertensie(?) Hypercholesterolemie (?)

Anamnesis Om 18.30 pijn op de borst met uitstraling naar de linkermoe, zweten, misselijk. Ambulance gebeld en bij aansluiten monitor beeld van acuut ondervoorhoede achterstand. 500 mg aspiric iv, ticagrelor 180 mg oraal, heparine, zofran eenmalig, 3x NTG spray, HD stabiel gebleven. Medicatie bij presentatie: Geen.

Lichamelijk onderzoek. Gauw, vegetatief, Halsvenen niet gestud. Cor s1 s2 geen souffles. Pulm schoon. Extre warm en slank.

Aanvullend onderzoek. AMBU ECG: Sinusritme, STEMI inferior III/II Cvermoedelijk RCA.

Coronair angiografie (...). Conclusie angio: 1-vatslijden...PCI

Conclusie en beleid

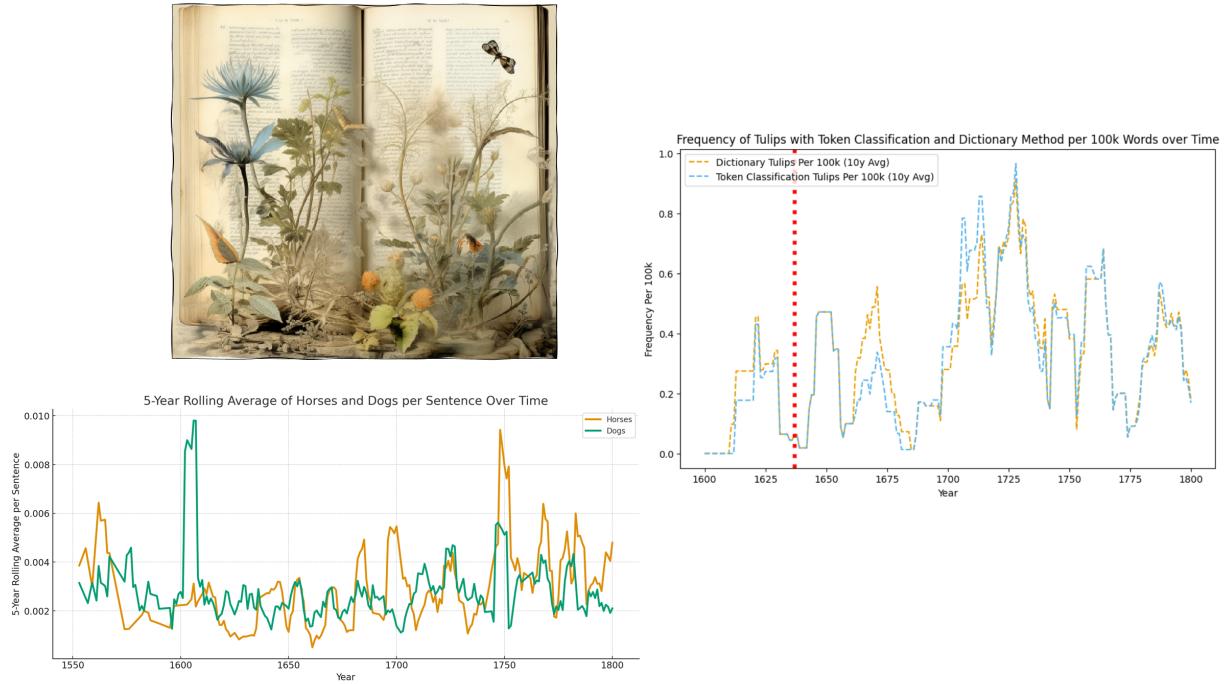
Bovengenoemde <LEEFIJD-1> jarige man, blanco cardiale voorgeschiedenis, werd gepresenteerd vanwege een STEMI inferior waarvoor een spoed PCI werd verricht van de mid-RCA. Er bestaan geen relevante neverletsel. Hij kon na de procedure worden overgeplaatst naar de CCU van het <INSTELLING-2>... Dank voor de snelle overname... Medicatie bij overplaatsing. Acetylsalicyzuur dispersablet 80 mg: oraal: 1x per dag 80 milligram; <DATUM-1> Ticagrelor tablet 90 mg: oraal: 2x per dag 90 milligram; <DATUM-1> Metoprolol tablet 50 mg: oraal: 2x per dag 25 milligram; <DATUM-1> Atorvastatine tablet 40 mg (als ca-zout-3-water); oraal: 1x per dag 40 milligram;

Samenvatting

Hoofddiagnose: STEMI inferior wv PCI RCA. Geen neverletsel. Nevendiagnoses: geen.

Complicates: geen Ontslag naar: CCU <INSTELLING-2>.

Opinions on Plants and Animals through Time Van Dalfsen et al. 2024



Course Logistics

Course materials

You can access the course materials quickly from

https://ayoubbagheri.nl/applied_tm/

Teachers




Pablo



Arjan



Jelte



Dong



Ayoub

Program

Time	Monday	Tuesday	Wednesday	Thursday	Friday
9:00 - 10:30	Lecture 1	Lecture 3	Lecture 5	Lecture 7	Lecture 9
	Break	Break	Break	Break	Break
10:50 – 12:00	Practical 1	Practical 3	Practical 5	Practical 7	Practical 9
12:00 – 12:30	Discussion 1	Discussion 3	Discussion 5	Discussion 7	Discussion 9
	Lunch	Lunch	Lunch	Lunch	Lunch
14:00 – 15:20	Lecture 2	Lecture 4	Lecture 6	Lecture 8	Lecture 10
	Break	Break	Break	Break	Break
15:30 – 16:30	Practical 2	Practical 4	Practical 6	Practical 8	Practical 10
16:30 – 17:00	Discussion 2	Discussion 4	Discussion 6	Discussion 8	Discussion 10

Goal of the course

- Text data are everywhere!
- A lot of world's data are in the format of unstructured text
- This course teaches
 - text mining techniques
 - using Python
 - on a variety of applications
 - in many domains.

Python?

How familiar are you with Python?

- What is your experience level with Python?



1 Go to wooclap.com
2 Enter the event code in the top banner

Event code
BUJKGM

A yellow rectangular box containing two steps for taking a poll on wooclap.com. Step 1 is to go to the website and step 2 is to enter the event code. The event code itself is displayed prominently in a large, bold, black font.

Python IDE?

- Which Python IDE do you mostly use? If you use more than one environment fill in the other text boxes.



1 Go to wooclap.com
2 Enter the event code in the top banner

Event code
BUJKGM

A yellow rectangular box containing two steps for taking a poll on wooclap.com. Step 1 is to go to the website and step 2 is to enter the event code. The event code itself is displayed prominently in a large, bold, black font.

Google Colab?

- How familiar are you with Google Colab? (1: limited to 5: expert)



Python

- Latest: Python 3.12.4
- Follow the tutorial on Python in Google Colab for the Applied Text Mining course: link
- Python For Beginners
 - <https://www.python.org/about/gettingstarted/>
- The Python Language Reference
 - <https://docs.python.org/3/reference/>
- Python 3.12.4 documentation
 - <https://www.python.org/doc/>

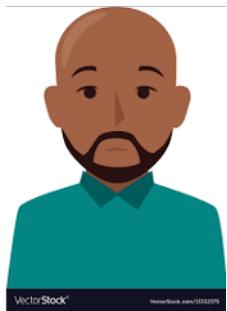
Google Colab

- Colaboratory, or “Colab” for short, allows you to write and execute Python in your browser, with
 - Zero configuration required
 - Free access to GPUs
 - Easy sharing
- [Intro](<https://colab.research.google.com/notebooks/intro.ipynb>)
- Cheat-sheet for Google Colab
- Keyboard shortcuts:

Actions	Colab	Jupyter
show keyboard shortcuts	Ctrl/Cmd M H	H
Insert code cell above	Ctrl/Cmd M A	A
Insert code cell below	Ctrl/Cmd M B	B
Delete cell/selection	Ctrl/Cmd M D	DD
Interrupt execution	Ctrl/Cmd M I	II
Convert to code cell	Ctrl/Cmd M Y	Y
Convert to text cell	Ctrl/Cmd M M	M
Split at cursor	Ctrl/Cmd M -	Ctrl Shift -

What is Text Mining?

Text mining in an example



- This is **Garry**!
- **Garry** works at Bol.com (a webshop in the Netherlands)
- He works in the dep of **Customer relationship management**.
- He uses Excel to read and search customers' reviews, extract aspects they wrote their reviews on, and identify their sentiments.
- Curious about his job? See two examples!

This is a nice book for both young and old. It gives beautiful life lessons in a fun way. Definitely worth the money!

+ Educational

+ Funny

+ Price

Nice story for older children.

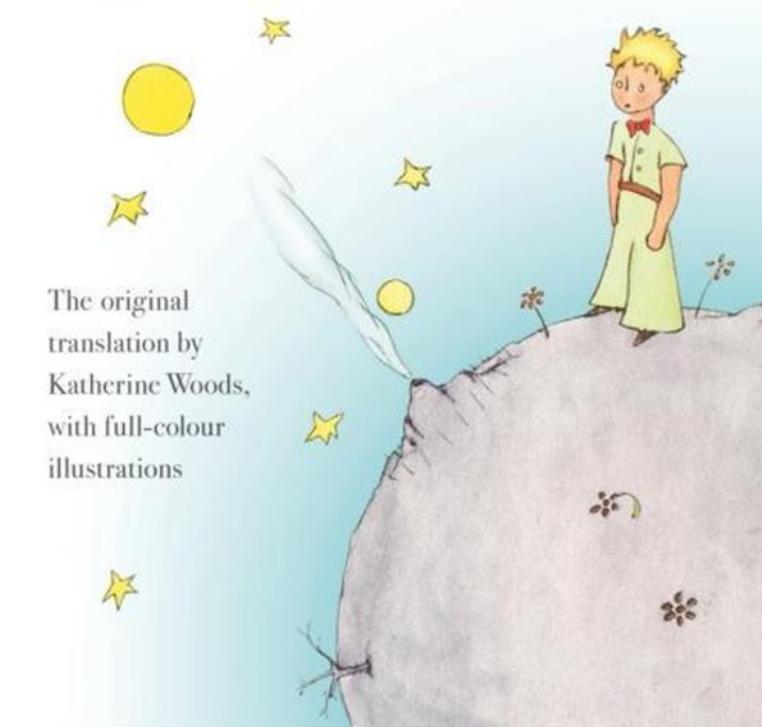
+ Funny

- Readability


ANTOINE DE SAINT-EXUPÉRY

The Little Prince

The original
translation by
Katherine Woods,
with full-colour
illustrations



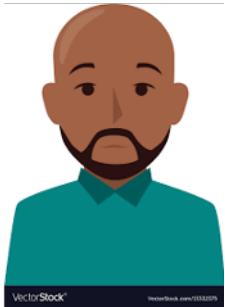
Example



- Garry likes his job a lot, but sometimes it is frustrating!
- This is mainly because their company is expanding quickly!
- Garry decides to hire **Larry** as his assistant.



Example



- Still, a lot to do for two people!
- Garry has some budget left to hire another assistant for couple of years!
- He decides to hire **Harry** too!
- Still, manual labeling using Excel is labor-intensive!



Language is hard!

- Different things can mean more or less the same (“data science” vs. “statistics”)
- Context dependency (“You have very nice shoes”);
- Same words with different meanings (“to sanction”, “bank”);
- Lexical ambiguity (“we saw her duck”)
- Irony, sarcasm (“That’s just what I needed today!”, “Great!”, “Well, what a surprise.”)
- Figurative language (“He has a heart of stone”)
- Negation (“not good” vs. “good”), spelling variations, jargon, abbreviations
- All the above are different over languages, 99% of work is on English!

Text mining

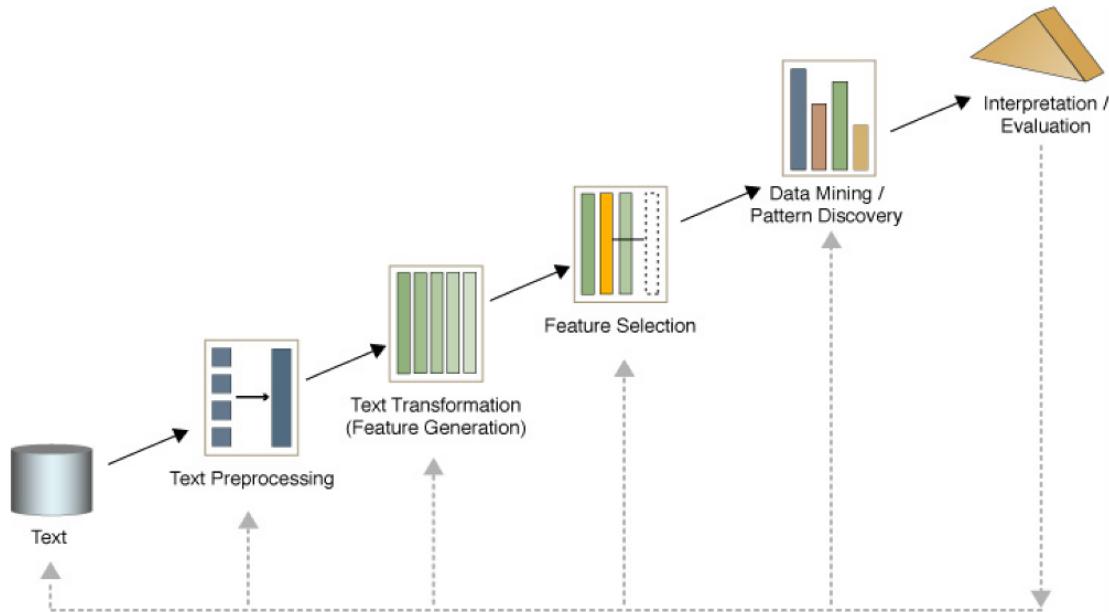
- “the discovery by computer of new, previously unknown information, by automatically extracting information from different written resources” Hearst (1999)
- Text mining is about looking for patterns in text, in a similar way that data mining can be loosely described as looking for patterns in data.
- Text mining describes a set of linguistic, statistical, and machine learning techniques that model and structure the information content of textual sources. (Wikipedia)

Can be quite effective!

- We won’t solve linguistics ...
- In spite of the problems, text mining can be quite effective!

Process & Tasks


Text mining process



Text mining tasks

- Text classification
- Text clustering
- Sentiment analysis
- Feature selection
- Topic modelling
- Responsible text mining
- Text summarization

And more in NLP



source: <https://nlp.stanford.edu/~wcmac/papers/20140716-UNLU.pdf>

Text Preprocessing

Text preprocessing

- is an approach for cleaning and noise removal of text data.
- brings your text into a form that is analyzable for your task.
- transforms text into a more digestible form so that machine learning algorithms can perform better.

Typical steps

- Tokenization (“text”, “ming”, “is”, “the”, “best” , “!”)
- Stemming (“lungs”→“lung”) or Lemmatization (“were”→“is”)
- Lowercasing (“Disease”→“disease”)
- Stopword removal (“text ming is best!”)
- Punctuation removal (“text ming is the best”)
- Number removal (“I42”→“I”)
- Spell correction (“hart”→“heart”)

Not all of these are appropriate at all times!

Tokenization/Segmentation

- Split text into words and sentences

There was an earthquake
near D.C. I've even felt it in
Philadelphia, New York,
etc.

There + was + an +
earthquake + near + D.C.

I + ve + even + felt + it + in +
Philadelphia, + New + York, +
etc.

N-grams

- N-grams: a contiguous sequence of N tokens from a given piece of text
 - E.g., ‘Text mining is to identify useful information.’
 - Bigrams: ‘text_mining’, ‘mining_is’, ‘is_to’, ‘to_identify’, ‘identify_useful’, ‘useful_information’, ‘information_’.
- Pros: capture local dependency and order
- Cons: increase the vocabulary size

Part Of Speech (POS) tagging

- Annotate each word in a sentence with a part-of-speech.

I ate the spaghetti with meatballs.
Pro V Det N Prep N

John saw the saw and decided to take it to the table.
PN V Det N Con V Part V Pro Prep Det N

- Useful for subsequent syntactic parsing and word sense disambiguation.

Vector Space Model

Basic idea

- Text is “unstructured data”
- How do we get to something structured that we can compute with?
- **Text must be represented somehow**
- Represent the text as something that makes sense to a computer

How to represent a document

- Represent by a string?
 - No semantic meaning
- Represent by a list of sentences?
 - Sentence is just like a short document (recursive definition)
- Represent by a vector?
 - A vector is an ordered finite list of numbers.

Vector space model

- A vector space is a collection of vectors
- Represent documents by concept vectors
 - Each concept defines one dimension
 - k concepts define a high-dimensional space
 - Element of vector corresponds to concept weight

Vector space model

- Distance between the vectors in this concept space
 - Relationship among documents
- The process of converting text into numbers is called Vectorization

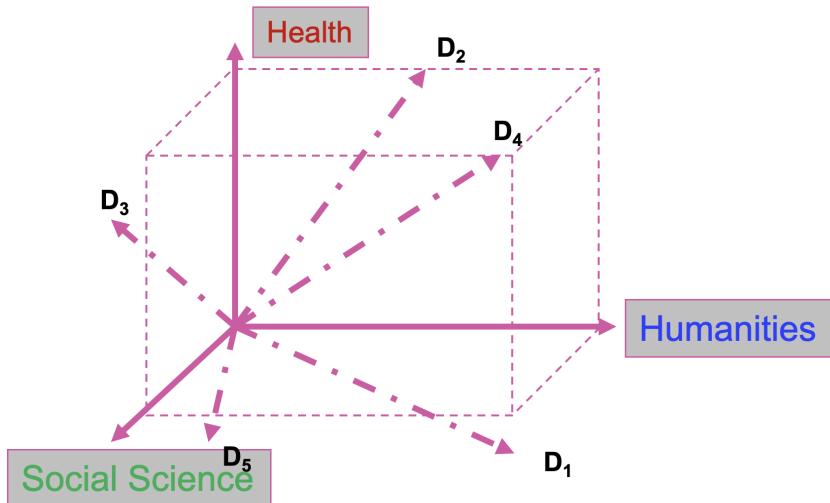
Vector space model

- Terms are generic features that can be extracted from text
- Typically, terms are single words, keywords, n-grams, or phrases
- Documents are represented as vectors of terms
- Each dimension (concept) corresponds to a separate term

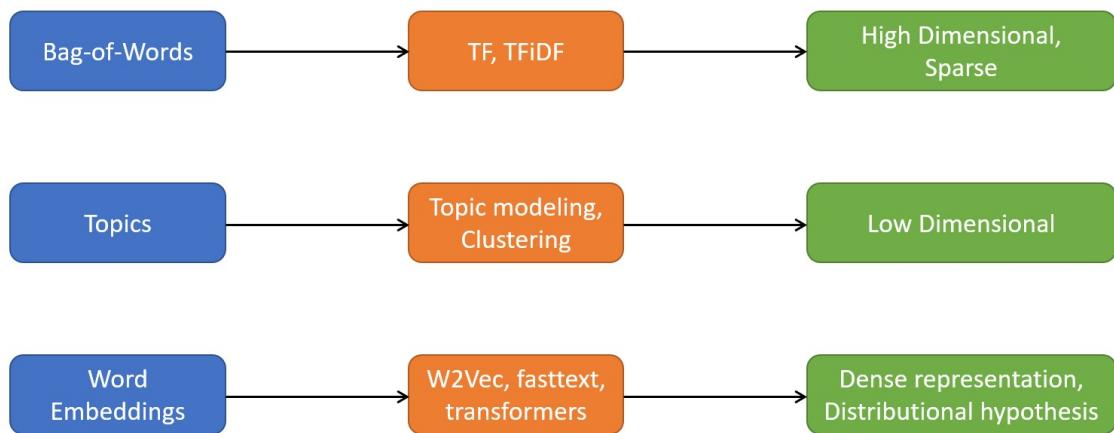
$$d = (w_1, \dots, w_n)$$

An illustration of VS model

- All documents are projected into this concept space



VSM: How do we represent vectors?



Bag of Words (BOW)

- *Terms* are words (more generally we can use n-grams)
- *Weights* are number of occurrences of the terms in the document
 - Binary
 - Term Frequency (TF)
 - Term Frequency inverse Document Frequency (TFiDF)

Binary

- Doc1: Text mining is to identify useful information.

- Doc2: Useful information is mined from text.
- Doc3: Apple is delicious.

	text	information	identify	mining	mined	is	useful	to	from	apple	delicious
Doc1	1	1	1	1	0	1	1	1	0	0	0
Doc2	1	1	0	0	1	1	1	0	1	0	0
Doc3	0	0	0	0	0	1	0	0	0	1	1

Term Frequency

- Idea: a term is more important if it occurs more frequently in a document
- TF formulas
 - Let $t(c, d)$ be the frequency count of term t in doc d
 - Raw TF: $tf(t, d) = c(t, d)$

TF: Document - Term Matrix (DTM)

Bag of words

- d1: "And God said, Let there be light: and there was light."
- d2: "And God saw the light, that it was good: and God divided the light from the darkness."
- d3: "And God called the light Day, and the darkness he called Night. And the evening and the morning were the first day."

"Document - Term matrix" (DTM) (raw word counts)

	light	god	darkness	called	day	let	said	divided	good	saw	evening	first	morning	night
d1	2	1	0	0	0	1	1	0	0	0	0	0	0	0
d2	2	2	1	0	0	0	0	1	1	1	0	0	0	0
d3	1	1	1	2	2	0	0	0	0	0	1	1	1	1

TFIDF

- Idea: a term is more discriminative if it occurs a lot but only in fewer documents

Let $n_{d,t}$ denote the number of times the t -th term appears in the d -th document.

$$TF_{d,t} = \frac{n_{d,t}}{\sum_i n_{d,i}}$$

Let N denote the number of documents and N_t denote the number of documents containing the t -th term.

$$IDF_t = \log\left(\frac{N}{N_t}\right)$$

TFIDF weight:

$$w_{d,t} = TF_{d,t} \cdot IDF_t$$

TFIDF: Document - Term matrix (DTM)

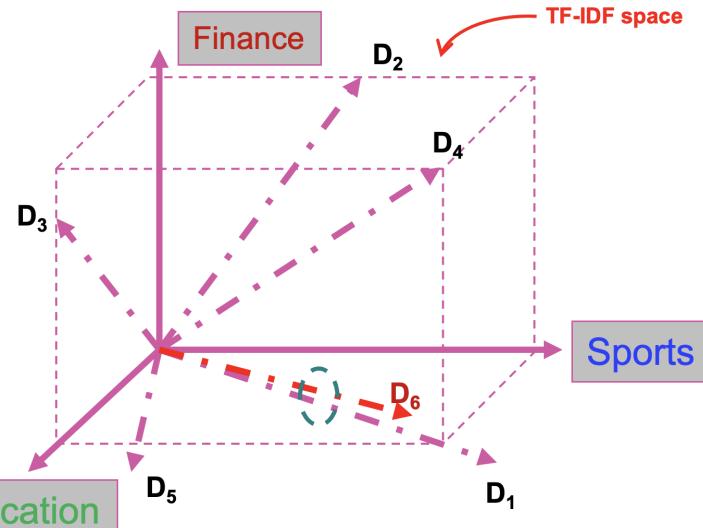
Bag of words

- d1: "And God said, Let there be light: and there was light."
- d2: "And God saw the light, that it was good: and God divided the light from the darkness."
- d3: "And God called the light Day, and the darkness he called Night. And the evening and the morning were the first day."

"Document - Term matrix" (DTM) (tf-idf)

	light	god	darkness	called	day	let	said	divided	good	saw	evening	first	morning	night
d1	0	0	0.000	0.0	0.0	1.1	1.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0
d2	0	0	0.405	0.0	0.0	0.0	0.0	1.1	1.1	1.1	0.0	0.0	0.0	0.0
d3	0	0	0.405	2.2	2.2	0.0	0.0	0.0	0.0	0.0	1.1	1.1	1.1	1.1

How to define a good similarity metric?



How to define a good similarity metric?

- Euclidean distance

$$dist(d_i, d_j) = \sqrt{\sum_{t \in V} [tf(t, d_i)idf(t) - tf(t, d_j)idf(t)]^2}$$

- Longer documents will be penalized by the extra words
- We care more about how these two vectors are overlapped

- Cosine similarity

- Angle between two vectors:

$$\text{cosine}(d_i, d_j) = \frac{V_{d_i}^T V_{d_j}}{|V_{d_i}|_2 \times |V_{d_j}|_2} \leftarrow \text{TF-IDF vector}$$

- Documents are normalized by length

Next

- Text classification

Summary

Summary

- Text data are everywhere!
- Language is hard!
- The basic problem of text mining is that text is not a neat data set
- Solution: text pre-processing & VSM

Practical 1