

Deep Learning for Text 1

Applied Text Mining

Ayoub Bagheri

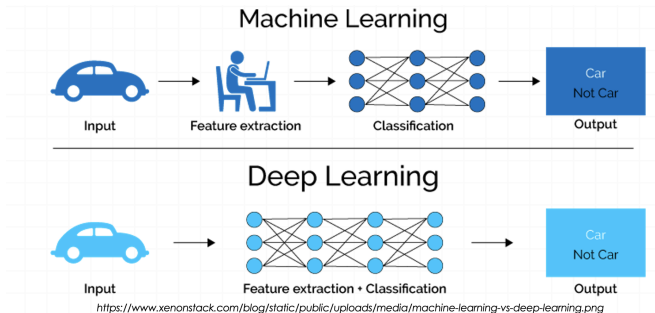
Lecture plan

1. Feed-forward neural networks
2. Recurrent neural networks
 - 2.1 SRN
 - 2.2 LSTM
 - 2.3 Bi-LSTM
 - 2.4 GRU

What is Deep Learning (DL) ?


A machine learning subfield of learning representations of data.
Exceptional effective at learning patterns.

Deep learning algorithms attempt to learn (multiple levels of) representation by using a hierarchy of multiple layers.




Feed-forward neural networks

- ▶ A typical multi-layer network consists of an input, hidden and output layer, each fully connected to the next, with activation feeding forward.




- ▶ The weights determine the function computed.

Feed-forward neural networks



Feed-forward neural networks



Weights

$$h = \sigma(W_1 x + b_1)$$
$$y = \sigma(W_2 h + b_2)$$

Activation functions

4 + 2 = 6 neurons (not counting inputs)
[3 x 4] + [4 x 2] = 20 weights
4 + 2 = 6 biases

26 learnable **parameters**

One forward pass

Text (input) representation

TFIDF

Word embeddings

....

0.2	-0.5	0.1
2.0	1.5	1.3
0.5	0.0	0.25
-0.3	2.0	0.0

W

0.1
0.2
0.3

x_i

+

1.0
3.0
0.025
0.0

b

=

0.95
3.89
0.15
0.37

$\sigma(x_i; W, b)$

very positive

positive

negative

very negative

Training

<https://medium.com/@ramrajchandradevan/the-evolution-of-gradient-descend-optimization-algorithm-4106a6702d39>

Optimize objective/cost function $J(\theta)$

Generate error signal that measures difference between predictions and target values

Use error signal to change the weights and get more accurate predictions

Subtracting a fraction of the gradient moves you towards the (local) minimum of the cost function

Updating weights


objective/cost function $J(\theta)$

Update each element of θ :

$$\theta_j^{new} = \theta_j^{old} - \alpha \frac{d}{d\theta_j^{old}} J(\theta)$$

Matrix notation for all parameters (α : learning rate):

$$\theta_j^{new} = \theta_j^{old} - \alpha \nabla_{\theta} J(\theta)$$




Recursively apply chain rule though each node

Non-linearities, old and new


logistic ("sigmoid")

$$f(z) = \frac{1}{1 + \exp(-z)}$$




tanh

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$




hard tanh

$$\text{HardTanh}(z) = \begin{cases} -1 & \text{if } z < -1 \\ z & \text{if } -1 \leq z \leq 1 \\ 1 & \text{if } z > 1 \end{cases}$$




(Rectified Linear Unit)
ReLU

$$\text{ReLU}(z) = \max(z, 0)$$



Leaky ReLU /
Parametric ReLU



tanh is just a rescaled and shifted sigmoid (2 × as steep, [-1,1]):


$$\tanh(z) = 2\text{logistic}(2z) - 1$$

Logistic and tanh are still used (e.g., logistic to get a probability)

However, now, for deep networks, the first thing to try is ReLU: it trains quickly and performs well due to good gradient backflow.

ReLU has a negative "dead zone" that recent proposals mitigate
GELU is frequently used with Transformers (BERT, RoBERTa, etc.)

Swish [arXiv:1710.05941](https://arxiv.org/abs/1710.05941)
 $\text{swish}(x) = x \cdot \text{logistic}(x)$





GELU [arXiv:1606.08415](https://arxiv.org/abs/1606.08415)
 $\text{GELU}(x)$
 $= x \cdot P(X \leq x), X \sim N(0,1)$
 $\approx x \cdot \text{logistic}(1.702x)$



Non-linearities (i.e., “ f ” on previous slide): Why they’re needed

- Neural networks do function approximation, e.g., regression or classification
 - Without non-linearities, deep neural networks can’t do anything more than a linear transform
 - Extra layers could just be compiled down into a single linear transform: $W_1 W_2 x = Wx$
 - But, with more layers that include non-linearities, they can approximate any complex function!



Notes on training


- ▶ Not guaranteed to converge to zero training error, may converge to local optima or oscillate indefinitely.
- ▶ However, in practice, does converge to low error for many large networks on real data.
- ▶ Many epochs (thousands) may be required, hours or days of training for large networks.
- ▶ To avoid local-minima problems, run several trials starting with different random weights (*random restarts*).
 - ▶ Take results of trial with lowest training set error.
 - ▶ Build a committee of results from multiple trials (possibly weighting votes by training set accuracy).

Hidden unit representations

- ▶ Trained hidden units can be seen as newly constructed features that make the target concept linearly separable in the transformed space.
- ▶ On many real domains, hidden units can be interpreted as representing meaningful features such as vowel detectors or edge detectors, etc..
- ▶ However, the hidden layer can also become a distributed representation of the input in which each individual unit is not easily interpretable as a meaningful feature.

Overfitting

Learned hypothesis may fit the training data very well, even outliers (noise) but fail to generalize to new examples (test data)




<http://wiki.bethanycrane.com/overfitting-of-data>

error



Overfitting prevention

- ▶ Running too many epochs can result in over-fitting.



- ▶ Keep a hold-out validation set and test accuracy on it after every epoch. Stop training when additional epochs actually increase validation error.
- ▶ To avoid losing training data for validation:
 - ▶ Use internal K-fold CV on the training set to compute the average number of epochs that maximizes generalization accuracy.
 - ▶ Train final network on complete training set for this many epochs.


Regularization

Dropout

Randomly drop units (along with their connections) during training

Each unit retained with fixed probability p , independent of other units


Hyper-parameter p to be chosen (tuned)



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting." Journal of machine learning research (2014)

Determining the best number of hidden units

- ▶ Too few hidden units prevents the network from adequately fitting the data.
- ▶ Too many hidden units can result in over-fitting.



- ▶ Use internal cross-validation to empirically determine an optimal number of hidden units.
- ▶ Hyperparameter tuning



Recurrent Neural Networks

Recurrent Neural Network (RNN)

- ▶ Add feedback loops where some units' current outputs determine some future network inputs.
- ▶ RNNs can model dynamic finite-state machines, beyond the static combinatorial circuits modeled by feed-forward networks.

Simple Recurrent Network (SRN)


- ▶ Initially developed by Jeff Elman (“*Finding structure in time*,” 1990).
- ▶ Additional input to hidden layer is the state of the hidden layer in the previous time step.



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>


Unrolled RNN

- Behavior of RNN is perhaps best viewed by “unrolling” the network over time.



Training RNNs

- ▶ RNNs can be trained using “backpropagation through time.”
- ▶ Can viewed as applying normal backprop to the unrolled network.



LSTM

Vanishing gradient problem

Suppose we had the following scenario:

Day 1: Lift Weights

Day 2: Swimming

Day 3: At this point, our model must decide whether we should take a rest day or yoga. Unfortunately, it only has access to the previous day. In other words, it knows we swam yesterday but it doesn't know whether had taken a break the day before.

Therefore, it can end up predicting yoga.

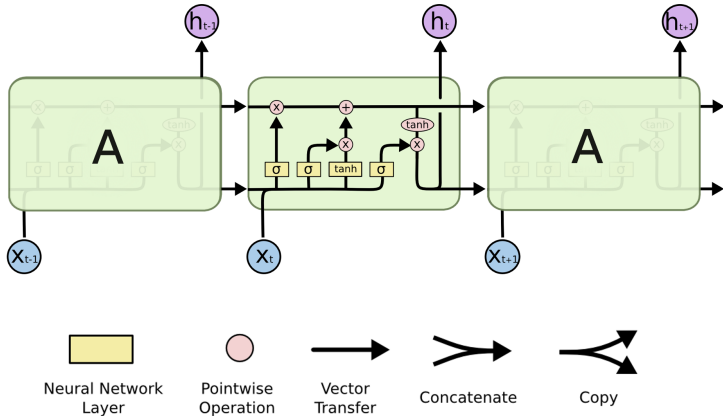
- ▶ Backpropagated errors multiply at each layer, resulting in exponential decay (if derivative is small) or growth (if derivative is large).
- ▶ Makes it very difficult train deep networks, or simple recurrent networks over many time steps.
- ▶ LSTMs were invented, to get around this problem.

<https://towardsdatascience.com/>

Long Short Term Memory

- ▶ LSTM networks, add additional gating units in each memory cell.
 - ▶ Forget gate
 - ▶ Input gate
 - ▶ Output gate
- ▶ Prevents vanishing/exploding gradient problem and allows network to retain state information over longer periods of time.

LSTM network architecture | <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>




->

-> -> ->

->


Bi-directional LSTM (Bi-LSTM)

- ▶ Separate LSTMs process sequence forward and backward and hidden layers at each time step are concatenated to form the cell output.



Gated Recurrent Unit (GRU)

- ▶ Alternative RNN to LSTM that uses fewer gates (Cho, et al., 2014)
 - ▶ Combines forget and input gates into “update” gate.
 - ▶ Eliminates cell state vector



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Attention

- ▶ For many applications, it helps to add “attention” to RNNs.
- ▶ Allows network to learn to attend to different parts of the input at different time steps, shifting its attention to focus on different aspects during its processing.
- ▶ Used in image captioning to focus on different parts of an image when generating different parts of the output sentence.
- ▶ In MT, allows focusing attention on different parts of the source sentence when generating different parts of the translation.

Summary

Summary

- ▶ Deep learning can be applied for automatic feature engineering
- ▶ Recurrent neural networks are ideal for sequential data such as text

Practical 6